

Lab-1 Report

Original Labs: [AI-Systems](#)

Experimental Enviroments

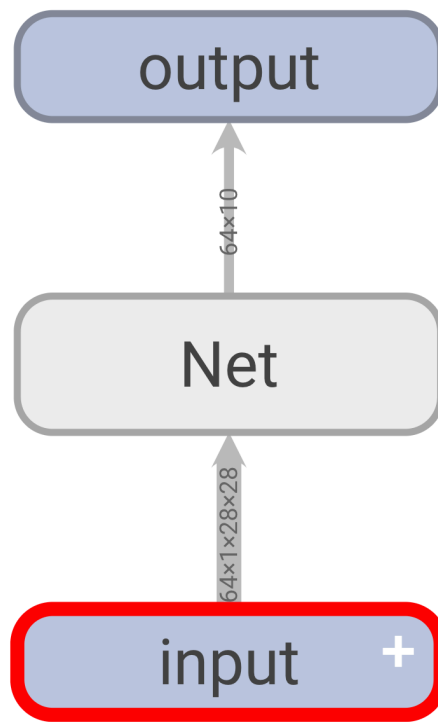
Hardware	CPU (# vCPU)	2
	GPU	/
Software	OS Version	macOS Mojave
	Package	PyTorch-1.9.0
	CUDA Version	/

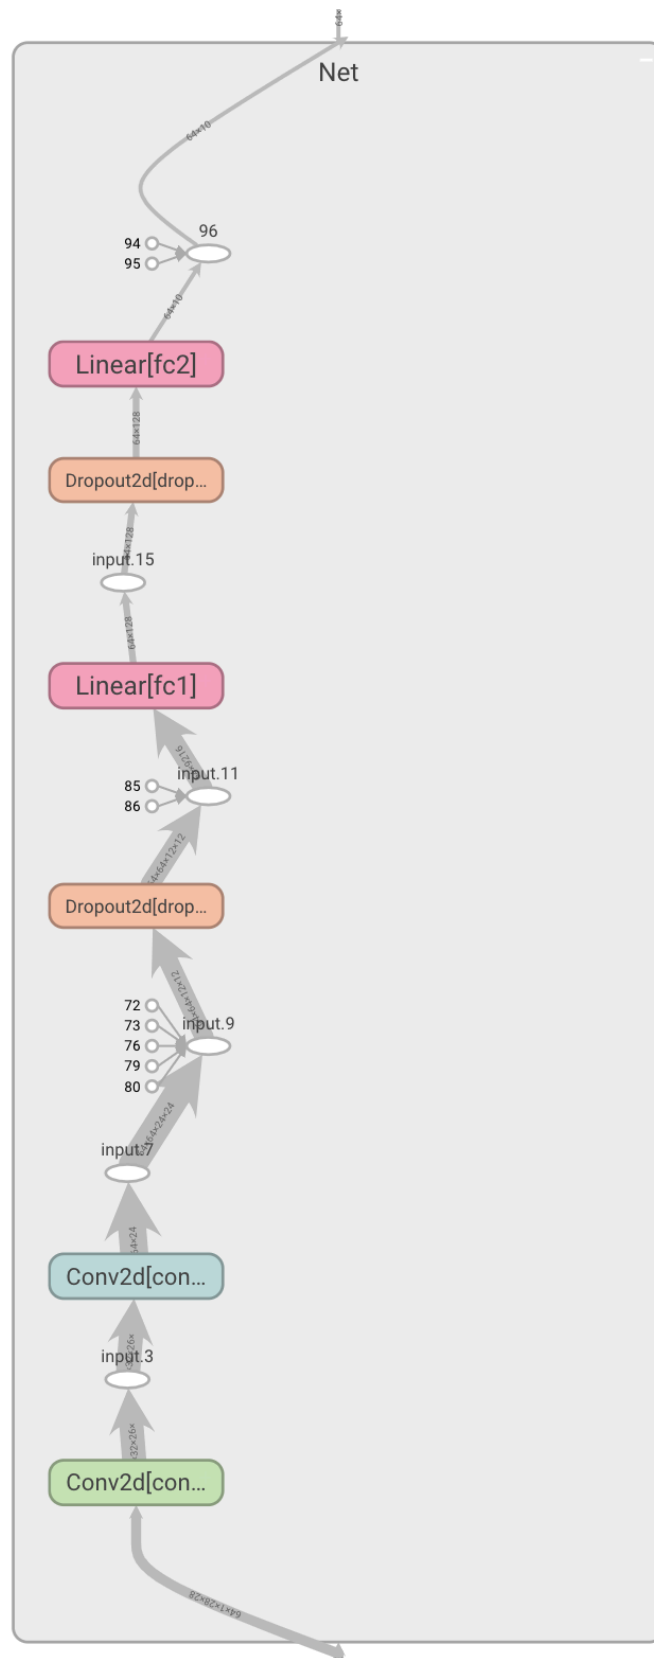
Experimental Results

Visualizing Models and Results

Note: use `tensorboard --logdir=logs`

Network Data Trend






Loss and Accuracy Trend

Accuracy




Accuracy/train
tag: Accuracy/train



Loss



Loss/train
tag: Loss/train



Operations

from profile results:

-----			-----	-----	-----	-----
-- -----						
			Name	Self CPU %	Self CPU	CPU total %
total	CPU time avg	# of Calls				CPU
-----				-----	-----	-----
-- -----						
			aten::conv2d	0.71%	903.317us	32.73%
41.396ms	41.396ms	1				
			aten::convolution	0.01%	17.259us	32.02%
40.492ms	40.492ms	1				
			aten::_convolution	0.61%	775.225us	32.01%
40.475ms	40.475ms	1				
			aten::mkldnn_convolution	31.36%	39.654ms	31.39%
39.700ms	39.700ms	1				
			aten::empty	0.02%	25.299us	0.02%
25.299us	25.299us	1				
			aten::as_strided_	0.02%	20.716us	0.02%
20.716us	20.716us	1				
			aten::relu	0.47%	597.455us	1.53%
1.940ms	1.940ms	1				
			aten::clamp_min	0.02%	27.279us	1.06%
1.343ms	1.343ms	1				
			aten::empty	0.00%	5.226us	0.00%
5.226us	5.226us	1				
			aten::clamp_min	1.04%	1.310ms	1.04%
1.310ms	1.310ms	1				
			aten::conv2d	0.00%	4.681us	35.13%
44.421ms	44.421ms	1				
			aten::convolution	0.00%	4.980us	35.12%
44.416ms	44.416ms	1				
			aten::_convolution	0.01%	14.663us	35.12%
44.411ms	44.411ms	1				
			aten::mkldnn_convolution	35.09%	44.379ms	35.11%
44.397ms	44.397ms	1				
			aten::empty	0.01%	14.540us	0.01%
14.540us	14.540us	1				
			aten::as_strided_	0.00%	3.650us	0.00%
3.650us	3.650us	1				
			aten::relu	0.01%	15.720us	0.94%
1.183ms	1.183ms	1				
			aten::clamp_min	0.00%	5.479us	0.92%
1.167ms	1.167ms	1				
			aten::empty	0.00%	2.878us	0.00%
2.878us	2.878us	1				

	aten::clamp_min	0.92%	1.159ms	0.92%
1.159ms	1.159ms 1			
	aten::max_pool2d	1.00%	1.266ms	10.86%
13.731ms	13.731ms 1			
	aten::max_pool2d_with_indices	9.86%	12.466ms	9.86%
12.466ms	12.466ms 1			
	aten::feature_dropout	0.42%	533.460us	4.71%
5.959ms	5.959ms 1			
	aten::empty	0.01%	9.312us	0.01%
9.312us	9.312us 1			
	aten::bernoulli_	2.31%	2.918ms	2.34%
2.959ms	2.959ms 1			
	aten::empty	0.03%	40.485us	0.03%
40.485us	40.485us 1			
	aten::div_	0.04%	54.159us	0.56%
704.197us	704.197us 1			
	aten::to	0.48%	603.033us	0.51%
650.038us	650.038us 1			
	aten::empty_strided	0.00%	3.819us	0.00%
3.819us	3.819us 1			
	aten::copy_	0.03%	43.186us	0.03%
43.186us	43.186us 1			
	aten::mul	1.39%	1.753ms	1.39%
1.753ms	1.753ms 1			
	aten::flatten	0.02%	19.643us	0.03%
41.842us	41.842us 1			
	aten::view	0.02%	22.199us	0.02%
22.199us	22.199us 1			
	aten::linear	2.11%	2.665ms	9.68%
12.240ms	12.240ms 1			
	aten::t	0.48%	609.637us	0.60%
757.760us	757.760us 1			
	aten::transpose	0.11%	140.705us	0.12%
148.123us	148.123us 1			
	aten::as_strided	0.01%	7.418us	0.01%
7.418us	7.418us 1			
	aten::addmm	6.93%	8.767ms	6.97%
8.817ms	8.817ms 1			
	aten::expand	0.02%	24.723us	0.02%
29.151us	29.151us 1			
	aten::as_strided	0.00%	4.428us	0.00%
4.428us	4.428us 1			
	aten::copy_	0.02%	21.343us	0.02%
21.343us	21.343us 1			
	aten::relu	0.01%	12.372us	0.04%
46.287us	46.287us 1			
	aten::clamp_min	0.01%	14.627us	0.03%
33.915us	33.915us 1			

	aten::empty	0.00%	4.259us	0.00%
4.259us	4.259us	1		
	aten::clamp_min	0.01%	15.029us	0.01%
15.029us	15.029us	1		
	aten::feature_dropout	0.01%	16.409us	0.18%
225.300us	225.300us	1		
	aten::empty	0.00%	2.341us	0.00%
2.341us	2.341us	1		
	aten::bernoulli_	0.09%	116.905us	0.09%
118.090us	118.090us	1		
	aten::empty	0.00%	1.185us	0.00%
1.185us	1.185us	1		
	aten::div_	0.02%	22.843us	0.03%
35.798us	35.798us	1		
	aten::to	0.00%	4.860us	0.01%
12.955us	12.955us	1		
	aten::empty_strided	0.00%	2.912us	0.00%
2.912us	2.912us	1		
	aten::copy_	0.00%	5.183us	0.00%
5.183us	5.183us	1		
	aten::mul	0.04%	52.662us	0.04%
52.662us	52.662us	1		
	aten::linear	0.00%	5.098us	0.08%
96.324us	96.324us	1		
	aten::t	0.01%	9.217us	0.01%
17.823us	17.823us	1		
	aten::transpose	0.00%	5.864us	0.01%
8.606us	8.606us	1		
	aten::as_strided	0.00%	2.742us	0.00%
2.742us	2.742us	1		
	aten::addmm	0.05%	62.877us	0.06%
73.403us	73.403us	1		
	aten::expand	0.00%	2.922us	0.00%
3.957us	3.957us	1		
	aten::as_strided	0.00%	1.035us	0.00%
1.035us	1.035us	1		
	aten::copy_	0.01%	6.569us	0.01%
6.569us	6.569us	1		
	aten::log_softmax	0.51%	642.944us	4.10%
5.185ms	5.185ms	1		
	aten::_log_softmax	3.58%	4.531ms	3.59%
4.542ms	4.542ms	1		
	aten::empty	0.01%	10.993us	0.01%
10.993us	10.993us	1		

Self CPU time total: 126.464ms

Compare Different Batch Size

Training 1 epochs:

Batch Size	Comparison
1	Self CPU time total: 51.256ms (51.256 ms/batch) Test set: Average loss: 0.1159, Accuracy: 9749/10000 (97%)
16	52.227ms (3.2641875ms/batch) Test set: Average loss: 0.0586, Accuracy: 9829/10000 (98%)
64	CPU time total: 126.464ms (1.976 ms/batch) Test set: Average loss: 0.0677, Accuracy: 9782/10000 (98%)